

# Supplementary Information

## AIRRRPrep: Workflows for Preparing Bulk and Single-Cell AIRR-Seq Data for Annotation

Parsa Esmaeili, Fatemeh Chaker Hosseini Zavareh, Nika Abdollahi and Changiz Eslahchi

### Overview of the Supplementary Information

This supplementary information supports the description, validation and benchmarking of AIRRRPrep presented in the main text, and is organized so that each item can be located by what it is needed for. Supplementary Figures S1–S4 are representative screenshots of the application: an example single-cell workflow diagram, a custom paired-read bulk workflow diagram, a stepwise read-retention plot, and a summary of which operations may act on Read 1, Read 2, or both. Supplementary Table S1 states, for bulk and single-cell data, what file types AIRRRPrep accepts as input and what it returns as output, and lists the seven single-cell input specifications. Supplementary Table S2 lists the 47 bulk operations AIRRRPrep exposes across 13 pRESTO tools, states the scope of their component-level comparison with native pRESTO, and reports the separate USEARCH-specific test that is not part of that comparison. Supplementary Note S1 explains how AIRRRPrep checks, both before a job is submitted and while a custom workflow is being assembled, that a chosen sequence of operations is compatible with the files supplied, and reports the positive and negative tests of that check. Supplementary Note S2 describes how a browser session and its uploaded files are managed, how a workflow can be saved, exported and re-imported, and how long-running jobs are tracked, notified and recovered. Supplementary Tables S3–S5 report, for each of the three predefined bulk workflows, the number of records retained after every processing step and the record-level comparison with the equivalent command-line pRESTO scripts. Supplementary Table S6 compares AIRRRPrep’s harmonized single-cell output record by record with three Cell Ranger representations of a public 10x Genomics dataset, including the join between the harmonized output and the lossless record of the annotations it omits, and Supplementary Table S7 reports the remaining single-cell input specifications: generic receptor FASTA (including inputs that must be refused) and the TRUST4-based routes. Supplementary Table S8 compares execution time between AIRRRPrep and command-line pRESTO for the three predefined workflows, first in aggregate and then broken down by individual processing step, and Supplementary Table S9 lists the parameters and the full reference command-line scripts used for that comparison, so that it can be reproduced exactly. Supplementary Table S10 lists the software and hardware environment in which the benchmarks were run, and Supplementary Table S11 compares AIRRRPrep with pRESTO, nf-core/airrflow, VDJServer and ARGalaxy, with a source for each entry.

### Supplementary Figure S1. Single-cell workflow visualization example




Figure S1: Representative AIRRPrep single-cell workflow visualization for a Cell Ranger V(D)J filtered contigs input. The application displays the processing route associated with the selected single-cell input and allows the workflow graphic to be exported as a PNG; this figure is that export. The annotation fields the *Strip annotation fields* step removes from the harmonized record are not discarded: they are kept unchanged in `source_annotations.tsv`, keyed on the same `sequence_id` as `final.fasta` and `metadata.tsv`. The run therefore writes three record-level outputs — `final.fasta`, `metadata.tsv` and `source_annotations.tsv` — plus the run-level `provenance.json`, which the diagram does not show because it is written once for the run rather than per record (Supplementary Table S1).

## Supplementary Figure S2. Visualization of a custom paired-read bulk workflow




Figure S2: Representative AIRRPrep paired-read bulk workflow visualization. Read 1 (R1) and Read 2 (R2) are shown through the operations applied to each read stream and, where applicable, their transition to a merged sequence stream after assembly. The workflow graphic can be exported from AIRRPrep as a PNG.

## Supplementary Figure S3. Stepwise read-retention visualization example




Figure S3: AIRRPrep read-retention visualization for the UMI-barcoded 5' RACE MiSeq 325+275 workflow of Supplementary Table S3 (SRR4026043, first 25,000 read pairs). Each box names its operation and shows the records entering and retained; R1 and R2 are counted separately until assembly, after which the merged stream is followed. The visualization can be exported as a PNG.

## Supplementary Figure S4. Bulk read-stream behaviour and output rules




Figure S4: Read-stream behaviour of bulk operations. An operation may act on the Read 1 (R1) or Read 2 (R2) stream independently, require both streams and keep them separate (for example `PairSeq.py`), or merge the two streams into a single sequence stream (`AssemblePairs.py`). Read streams are called R1 and R2 throughout this work, never “lane”: a sequencing flow-cell lane is a different thing, and the word is used in this supplement only where a physical lane is meant (lane L001 in Supplementary Table S7). The four assembly operations, written `AssemblePairs.py` `align`, `join`, `reference` and `sequential` in this supplement and in every table, are the components the interface labels `AssembleSeq.align`, `AssembleSeq.join`, `AssembleSeq.reference` and `AssembleSeq.sequential`; that family is the only name that differs between the interface and pRESTO (Supplementary Table S2).

## Supplementary Table S1. Input and output by data category

| Data category    | Input                                                                                                            | Output                                                                                                                             |
|------------------|------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| Bulk AIRR-seq    | One sequence file, or matched R1 and R2 files of a paired-end run; FASTA or FASTQ, both files in the same format | FASTA, FASTQ or TSV                                                                                                                |
| Single-cell AIRR | One of the seven input specifications below                                                                      | <code>final.fasta</code> ,<br><code>metadata.tsv</code> ,<br><code>source_annotations.tsv</code> ,<br><code>provenance.json</code> |

| ID | Specification                 | Required files                                                                                                                                                              |
|----|-------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1  | Cell Ranger, all contigs      | <code>all_contig.fasta</code> + <code>all_contig_annotations.csv</code>                                                                                                     |
| 2  | Cell Ranger, filtered contigs | <code>filtered_contig.fasta</code> + <code>filtered_contig_annotations.csv</code>                                                                                           |
| 3  | Cell Ranger AIRR table        | <code>airr_rearrangement.tsv</code>                                                                                                                                         |
| 4  | TRUST4 output                 | <code>*_annot.fa</code> + <code>*_barcode_report.tsv</code> , from one TRUST4 run in barcode mode                                                                           |
| 5  | Raw 10x FASTQ                 | R1 FASTQ (cell barcode + UMI) + R2 FASTQ (cDNA); assembled with TRUST4                                                                                                      |
| 6  | 10x BAM                       | genome-aligned BAM with CB/UB tags (e.g. <code>possorted_genome_bam.bam</code> ); assembled with TRUST4                                                                     |
| 7  | Generic FASTA                 | receptor FASTA + barcode mapping file ( <code>sequence_id</code> , <code>cell_id</code> , <code>umi_count</code> , <code>read_count</code> ) or a header regular expression |

*Supplementary Table S1.* Supported input and output file types for bulk and single-cell AIRR-seq data, and the seven single-cell input specifications. The two Cell Ranger contig sets are counted as separate specifications because they are different files with different content. When an uploaded file set matches a specification only in part, the application names the companion files needed to complete it. Single-cell input is not processed through a user-composed chain, because upstream software may already have reconstructed contigs; each specification has one adapter. The single-cell output is AIRRPrep’s internal harmonized representation, not an AIRR Rearrangement record, with eight fields: `sequence_id` (the source record’s own identifier, unchanged), `sequence`, `cell_id`, `umi_count`, `read_count`, `is_cell`, `high_confidence` and `productive`. A value the source does not carry is left empty rather than inferred, and `productive` is filled only when the user opts in, for Cell Ranger input.

Upstream V(D)J calls, junctions and CDR3 annotations are excluded from that harmonized record so that every sequence is re-annotated with one reference and one procedure. Excluding them from the analysis schema does not destroy them: every field the uploaded file carried for a record — the contig annotations CSV or AIRR TSV row, the barcode mapping file row, the FASTA header, TRUST4’s `annot.fa` header and its barcode-report and `barcode_airr` rows — is written verbatim to `source_annotations.tsv`, one row per output record, keyed on the same `sequence_id`. An inner join on that column is total and collision-free, so any omitted annotation can be restored without the original upload. `provenance.json` records the name, size and SHA-256 checksum of every input and of every output including the sidecar, which files the sidecar was read from, the adapter options, software versions and parse counts; a checksum of the source file therefore verifies the identity of that file, while the sidecar is what preserves its annotation content. The join is covered by an automated test (`backend/tests/test_source_annotations.py`), which asserts that every harmonized record has exactly one sidecar row, that no identifier is duplicated or lost on either side, and that the fields the harmonized output omits are present in the sidecar with the values the source gave.

*Note:* In the hosted deployment an uploaded file may be up to 200 MB and one session may hold up to 2 GB of uploads; uploaded files are retained for ten days and can be reused across runs within that window, and a job recovered with a tracking code remains available for seven days. All four limits are deployment settings (Supplementary Note S2).

## Supplementary Table S2. Bulk pRESTO operations exposed by AIRRPrep

AIRRPrep exposes 47 operations drawn from 13 pRESTO tools. Each operation listed below has exactly one archived test case, run within AIRRPrep and natively on the same input and the two outputs compared; those 47 cases are the denominator of every count reported in this section, and the scope of the comparison is stated in the subsection after the table. A separate, additional test covers the three **ClusterSets** operations with the proprietary USEARCH, which is not part of the 47 and is reported on its own at the end of this section.

| Tool                           | Description                                                          | Operations exposed                                                                                                                                                           |
|--------------------------------|----------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>AlignSets.py</code>      | Multiple-aligns sets of sequences sharing the same annotation        | <code>muscle</code> , <code>offset</code> , <code>table</code> (3)                                                                                                           |
| <code>AssemblePairs.py</code>  | Assembles paired-end reads into a single sequence                    | <code>align</code> , <code>join</code> , <code>reference</code> , <code>sequential</code> (4)                                                                                |
| <code>BuildConsensus.py</code> | Builds a consensus sequence for each UMI/barcode group               | single operation (1)                                                                                                                                                         |
| <code>ClusterSets.py</code>    | Clusters read groups                                                 | <code>all</code> , <code>barcode</code> , <code>set</code> (3) <sup>a</sup>                                                                                                  |
| <code>CollapseSeq.py</code>    | Removes duplicate sequences while recording duplicate counts         | single operation (1)                                                                                                                                                         |
| <code>ConvertHeaders.py</code> | Converts sequence headers into pRESTO format                         | <code>454</code> , <code>genbank</code> , <code>generic</code> , <code>illumina</code> , <code>imgt</code> , <code>migec</code> , <code>sra</code> (7)                       |
| <code>EstimateError.py</code>  | Estimates error rates for UMI data                                   | <code>barcode</code> , <code>set</code> (2)                                                                                                                                  |
| <code>FilterSeq.py</code>      | Removes or modifies low-quality reads                                | <code>length</code> , <code>quality</code> , <code>missing</code> , <code>repeats</code> , <code>trimqual</code> , <code>maskqual</code> (6)                                 |
| <code>MaskPrimers.py</code>    | Identifies and removes or masks primers while annotating the match   | <code>align</code> , <code>extract</code> , <code>score</code> (3)                                                                                                           |
| <code>PairSeq.py</code>        | Synchronizes paired R1/R2 files and copies annotations between mates | single operation (1)                                                                                                                                                         |
| <code>ParseHeaders.py</code>   | Manipulates sequence annotations and can export them as a table      | <code>add</code> , <code>collapse</code> , <code>copy</code> , <code>delete</code> , <code>expand</code> , <code>merge</code> , <code>rename</code> , <code>table</code> (8) |
| <code>SplitSeq.py</code>       | Splits, groups, sorts, samples and selects sequence files            | <code>count</code> , <code>group</code> , <code>sample</code> , <code>samplepair</code> , <code>select</code> , <code>sort</code> (6)                                        |
| <code>UnifyHeaders.py</code>   | Unifies annotation fields by a grouping scheme                       | <code>consensus</code> , <code>delete</code> (2)                                                                                                                             |
| <b>13 tools</b>                |                                                                      | <b>47 operations</b>                                                                                                                                                         |

<sup>a</sup> Counted once each in the 47 cases, with VSEARCH 2.31.0 as the clustering backend — the open tool the container image carries and the one a deployment uses unless it is changed. AIRRPrep issues the clustering command itself rather than through the pRESTO helper, whose USEARCH command line the tested USEARCH build rejects; the clustering parameters exposed to the user are unchanged. The USEARCH backend is covered by the separate test described at the end of this section.

*Supplementary Table S2.* The 47 bulk-processing operations exposed by AIRRPrep across the 13 underlying pRESTO tools that provide them, with a short description of each tool’s function. Each operation has exactly one archived test case comparing it with the equivalent native pRESTO command, and those 47 cases are the denominator used throughout this section. The per-case record is released with the source code as `paper/test/component_matrix/component_matrix.tsv` and `component_matrix.json`, one row per case with the operation, the denominator it belongs to, the host it ran on, the input file, the parameters, the native command, the AIRRPrep call, the comparator, the status and the SHA-256 of each side’s output set; the separate USEARCH test is released beside it as `usearch_clustersets.tsv` and `usearch_clustersets.json`. The results are summarized below.

*Note:* This supplement uses pRESTO’s program names throughout. The interface labels one family differently: `AssembleSeq.align`, `.join`, `.reference` and `.sequential` in the interface, in saved workflows and in job logs are `AssemblePairs.py align`, `join`, `reference` and `sequential`; every other component is named `Tool.subcommand` after its pRESTO program. `ParseLog.py`, which parses pRESTO log files after a run, is not exposed as a pipeline component, because AIRRPrep reports the per-step record counts directly (Supplementary Figure S3).

## Validation of AIRRPrep on pRESTO Components

**Scope of the component comparison.** Each of the 47 operations in Supplementary Table S2 has one archived test case, run through AIRRPrep’s own step executor (`run_pipeline_step`, the function a job calls for each step, with the parameters the web interface sends) and as the equivalent native pRESTO command on the same input, with the two outputs compared. The inputs are the first 8,000 read pairs of SRR1383456 and intermediates derived from them with native pRESTO, so no AIRRPrep output is ever the input of a test; the `ConvertHeaders` cases use 500 of those reads with headers rewritten into each vendor format, and reference-guided assembly uses 300 read pairs because it calls BLAST once per pair. The case list, parameters, commands and per-case results are archived with the source code (`paper/test/component_matrix/`), as is the harness that produces them.

Outputs are compared by content, not by file layout. Sequence outputs are compared as record sets — identifier, sequence, quality string and the annotation field/value set — with pass and failed records compared separately (for read-pair assembly the two sides group failed records differently, pRESTO writing one file per input and AIRRPrep one file per step, so those records are compared as one set per side); an operation that fans out into several files must match one native file per AIRRPrep file; table outputs are compared as sets of rows keyed by column name; random sampling, which pRESTO does not seed, is checked for equal size, absence of duplicates and membership of every sampled record in the input; and clustering is checked for identical records apart from the cluster label together with an identical partition of records into clusters.

**Result.** Of the 47 cases, 46 are comparable and all 46 produced identical output on both sides, over 123,016 compared sequence records in the record-set cases plus the fan-out, table, sampling and clustering cases. The 47th, `ConvertHeaders.py migecc`, is *not comparable*: native pRESTO refuses its own documented input, so there is no native output to compare against, and AIRRPrep’s run of it is reported as a native failure rather than as agreement (below). Every case therefore carries exactly one of two statuses, *identical* or *not comparable (native pRESTO failed)*, and their counts are 46 and 1.

Five of the 47 cases were run in a Linux container rather than on the benchmark host, because they cannot run on Windows at all: pRESTO passes those external tools a temporary file while Python still holds it open, which Windows does not permit, so *both* implementations fail and nothing can be compared. This affects the three `ClusterSets` operations, where VSEARCH reports that it cannot open the uc file, and the two reference-guided assembly operations, where `makeblastdb`

writes no database and every read pair is consequently reported as unassembled. The container is built from the same pRESTO version and the official VSEARCH and BLAST+ releases, pinned by checksum (Supplementary Table S10). Which host ran each case is recorded per case in the released report.

This per-operation matrix and the whole-workflow comparison of Supplementary Tables S3–S5 answer different questions. The workflows show that 11 operations behave identically in composition, over 25,000 read pairs and with each step’s output feeding the next; the matrix shows that each of the 46 comparable operations behaves identically in isolation, including those no predefined workflow uses. No claim of identical behaviour is made for `ConvertHeaders.py migec`, which has no native counterpart to be identical to.

**The one case that could not be compared.** `ConvertHeaders.py migec` is the 47th case, and native pRESTO refuses its own documented input. MIGEC writes every record as `@MIG UMI:barcode:count`, so the leading token is MIG in every record; pRESTO counts the records of an input file by building a Biopython index keyed on that token, which raises `Duplicate key 'MIG'` and ends the run before any header is converted. Any MIGEC file with more than one record is affected. AIRRPrep streams the file instead of indexing it and converts it with pRESTO’s own `convertMIGECHeader` function, so the conversion is pRESTO’s; what differs is only the file handling around it. Because native pRESTO writes no output, there is nothing to compare the 500 converted records against, and the case is reported as *not comparable*, not as agreement.

## Separate test: ClusterSets with USEARCH

The three `ClusterSets` cases of the matrix above use VSEARCH. USEARCH is tested separately, for two reasons: counting the same three operations twice would make the denominator ambiguous, and USEARCH is proprietary, so a reader cannot repeat that test without a licence of their own. It is therefore reported here, on its own, and is archived separately as `paper/test/component_matrix/usearch_clustersets.tsv` and `.json`, produced by `run_usearch_clustersets.py`.

Native pRESTO cannot run `ClusterSets.py` with USEARCH at all, because of the command line pRESTO constructs rather than any parameter chosen by the user. pRESTO builds its USEARCH invocation with the `-minwordmatches` option, which USEARCH v7.0.1090 – the build tested here, and the series pRESTO documents for this step – does not recognise; USEARCH exits with `Invalid command line, Unknown option minwordmatches`, and pRESTO reports the step as failed. Because the option is part of the command pRESTO constructs, the failure occurs for any input and cannot be avoided through the parameters the tool exposes. All three operations (`all`, `barcode`, `set`) fail identically.

AIRRPrep therefore issues the clustering command itself instead of delegating to the pRESTO helper. The USEARCH invocation omits the unsupported option; the VSEARCH invocation passes it in VSEARCH’s long-option form (`--minwordmatches`); and the thread count is omitted from the command when it has not been configured, rather than being interpolated as the literal string `None`. With that invocation USEARCH accepts the command line and starts, for all three operations.

The comparison cannot be carried through to an output check on the benchmark host, for the same reason the three VSEARCH cases of the matrix are run in the Linux container: pRESTO hands the external tool a temporary file that Python still holds open, and on Windows USEARCH then reports `Cannot open ...errno=13 Permission denied`. Doing for USEARCH what was done for VSEARCH would require a licensed Linux USEARCH binary inside the test container, which cannot be redistributed with the test suite. What is established here is therefore exactly this: the command line native pRESTO generates is rejected by USEARCH, and the one AIRRPrep

generates is not. The clustering parameters exposed to the user – identity threshold, length ratio, sequence window and choice of clustering tool – are unchanged either way.

USEARCH is proprietary and is not distributed with AIRRPrep or its container image; a deployment that holds a USEARCH licence supplies its own binary through the `USEARCH_BIN` setting, and selecting USEARCH without one fails with a message naming the licence requirement rather than a missing executable. No default selects it: clustering defaults to VSEARCH and reference-guided assembly to BLAST+, both of which the image carries. That the released tree and image contain no USEARCH binary, that nothing in the build installs one, and that no predefined workflow needs one, are asserted by automated tests (`backend/tests/test_no_proprietary_binaries.py`). The USEARCH and VSEARCH versions used are listed in Supplementary Table S10.

## Supplementary Note S1. Declarative Workflow Validation and Interactive Constraint Enforcement

Workflow validation in AIRRPrep is implemented within the framework itself and is not inherited from the underlying toolkits. pRESTO validates the arguments of an individual command at the time that command is executed, but it does not establish whether a sequence of commands is coherent with respect to the files actually supplied. In a command-line setting an incompatibility of this kind is therefore discovered only when the offending step runs, after the preceding steps have already been executed, and in some cases only in the form of a silently empty output file.

Every operation exposed by AIRRPrep declares a contract: the file types it accepts, the annotation capabilities it requires and the capabilities it produces, the companion files it needs (for example a primer FASTA), and whether it acts on one read stream, requires both the Read 1 (R1) and Read 2 (R2) streams, or merges them into one. Capabilities are tracked separately for the R1 and R2 streams. Before a job is dispatched, the selected workflow is replayed symbolically over the uploaded files: the file type and the capability state of each read stream are updated after every operation, and execution is refused at the first step whose contract is not satisfied. The resulting error identifies the failing step, the affected read stream and the specific unmet requirement. Representative constraints enforced by this mechanism include consensus building before barcode annotation has been created, paired operations on single-file input, quality-based filtering on FASTA input, which carries no quality scores, and primer masking without the required primer FASTA.

The same contracts are evaluated interactively while a custom workflow is being constructed. The set of selectable components is recomputed after every change to the uploaded input set or to the pipeline, so that components whose input types or required capabilities are not satisfied in the current state are disabled, whereas components become available as soon as their prerequisites are produced by an earlier step or supplied by an additional input file. Selecting an operation immediately updates which downstream operations are offered, uploading a different input type or a companion file re-enables the operations that depend on it, and a disabled component reports the requirement it is missing rather than being unavailable without explanation. Because only operations that are valid for the current data can be selected, assembling a valid pipeline does not require prior knowledge of the correct ordering of preprocessing steps. The same operation contracts are also used to generate the in-application documentation, so the documented requirements and the enforced requirements cannot diverge, and each predefined workflow is shown with a schematic of the library design it targets.

The same mechanism is applied a level deeper, to individual parameters rather than whole operations. Several operations take an annotation field name as a parameter – the field an operation groups, sorts or filters by, for example, or the annotations two paired-end reads exchange – and which field names are actually present depends on the specific workflow assembled so far: it begins with the fields detected in the uploaded file’s headers and changes as operations that add, rename, merge or remove annotations are added, so it cannot be read off the operation’s contract alone. AIRRPrep tracks this field-name set per read stream alongside the capability state described above, and wherever a parameter names an existing or new field, offers it as a choice drawn from that tracked set rather than requiring the name to be typed. When the uploaded header format has not yet been normalized to pRESTO’s convention – before the corresponding header-conversion operation has been applied, for instance – the tracked set may be incomplete, and a typed value remains available alongside it, since an unconfirmed field list must not silently reject a name it has simply not seen yet. The same field-existence check is applied once more, non-interactively, when the workflow is dispatched.

**Tests of the validator.** The job-launch validator is covered by an automated test suite (`backend/tests/test_workflow_validation.py`) that runs it on freshly written FASTQ and FASTA files without pRESTO annotations, the state of a new upload. The three predefined workflows must be accepted, and each of the following 14 deliberately invalid compositions must be refused at the step at fault:

| Class                      | Invalid composition (expected: refused)                                                                                                                                                                                        |
|----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Missing annotation         | consensus building on reads without a barcode; the UMI 2×250 workflow without the <code>PairSeq.py</code> step that copies the barcode to R2                                                                                   |
| Incorrect file type        | quality filtering on FASTA; R1 FASTQ with R2 FASTA                                                                                                                                                                             |
| Absent companion file      | primer masking without a primer FASTA; reference-guided assembly without a reference                                                                                                                                           |
| Invalid read-stream state  | <code>PairSeq.py</code> on single-file input; <code>AssemblePairs.py</code> on single-file input; an R2 operation on single-file input                                                                                         |
| Impossible operation order | an R2 operation after the streams were merged; <code>PairSeq.py</code> after assembly; any operation after a split into an unbounded number of files; a field parameter naming an annotation that does not exist at that point |
| Unknown operation          | a component name the server does not provide                                                                                                                                                                                   |

All 3 valid workflows were accepted and all 14 invalid compositions were refused with an error naming the failing step (no false acceptance and no false rejection on this set). The set is designed to cover each class of constraint the contracts express, not to estimate error rates over arbitrary user-built workflows.

## Supplementary Note S2. Managing Data, Workflows, and Running Jobs

**Sessions and uploaded files.** AIRRPrep requires no account. Work is organized in a browser session, which is issued its own secret token; every session, file and job route is authorized against that token, and only a hash of it is stored. Uploaded files belong to the session and are retained for ten days, so a second workflow can reuse a file that has already been uploaded rather than transferring it again; the interface shows the remaining retention time, so reuse is an informed choice. A session may hold up to 2 GB of uploads with a ceiling of 200 MB per file, and both limits, like the retention window, are deployment settings. Jobs previously run in a session remain listed with their status, a running job can be cancelled, and the output of an individual step – including the `_fail` file of a filtering step – can be downloaded on its own or the complete set of job outputs as a single archive. Because the service is public, uploads are restricted by file type and size, and requests are rate limited per client address.

**Saved and exported workflows.** A workflow assembled in the custom builder can be saved under the session and exported as a machine-readable definition, which can be imported again – into another session, or by another user – to reproduce the same operation sequence and parameters. A finished run can likewise be saved as a reusable workflow. Parameters that reference an uploaded file, such as a primer FASTA, are deliberately not embedded in the saved definition: uploads expire, and a stored reference to a file that no longer exists would fail at launch with an identifier the user cannot interpret. Such parameters are stripped on save and recorded instead as required inputs, each named together with the step that consumes it, so re-running asks for the primer FASTA again explicitly and by name. When the parameter named one of the session’s own uploads at the time of saving, the entry also records that file’s original name, size and SHA-256 checksum, and the definition records the AIRRPrep, pRESTO and Python versions under which it was saved and exported (an image digest as well, when the deployment sets it). The exported definition is therefore a reusable workflow definition that identifies every companion file by content; with the accessions of the sequencing inputs and the digest of the container image, it identifies the preprocessing that was applied.

**Long-running jobs, notification and recovery.** Before a job starts, AIRRPrep estimates its runtime from the size of the input and from the throughput of the operations the workflow contains, since throughput differs by more than an order of magnitude between, for example, a per-read quality predicate and a local alignment against every primer. Runs the estimator considers long may be started only from a session that has verified an email address; the estimate is a sizing heuristic and errs towards over-estimation, and the requirement is enforced by the backend rather than by the interface. Any session with a verified address is notified when its jobs start and when they finish or fail, whether or not the run was long.

The start notification carries a six-digit tracking code that restores the run on the page it was launched from, from any browser or machine, for seven days. A six-digit code is short enough to quote or write down, and correspondingly too small a space to serve as a credential on its own, so recovery requires the code together with the address it was issued to. For the same reason the link in the email pre-fills only the code, never the address.

**Handling of email addresses and recovery controls.** The address a user verifies is kept in plain text on the session record and on the records of jobs started from that session, because it is needed to deliver the start and completion notifications; those records expire with the session and job retention periods. Mail is sent by the backend over SMTP to the operator’s configured mail server. The tracking record that a recovery looks up does not contain the address: it stores an

HMAC-SHA256 of the normalized address, keyed with a server-side secret held in the deployment environment rather than in the data store, so a copy of the data store alone does not allow addresses to be recovered by hashing candidate addresses. The same code and a different address give the same failure as an unknown code, so the route does not reveal which codes exist. Verification codes expire after 15 minutes, are stored only as a salted hash, are invalidated after five wrong attempts, and cannot be resent to a session within 60 seconds. The recovery route requires a solved CAPTCHA before any lookup and is limited to ten requests per minute per client address (five per minute for sending a verification code); tracking codes expire seven days after their last use. All of these values are deployment settings.

Every statement in this paragraph is covered by an automated test (`backend/tests/test_security_controls.py`), which runs against an in-memory store with an injectable clock, so expiry and cooldowns are checked without a server and without waiting: code expiry and that a wrong guess does not extend it, the attempt limit and that it burns the code, the resend cooldown, storage as a salted hash only, address normalisation and the refusal of malformed addresses, the keyed HMAC on the tracking record, the identical answer for a wrong code and a wrong address, the absence of the address from the recovery link, the rate limits, and the consumption of a CAPTCHA on first use whether the answer was right or wrong. What the operator must supply rather than the code enforce — the HMAC secret, TLS in front of the service, the SMTP configuration, the retention policy of the plain-text addresses on session and job records — is set out in `SECURITY.md` in the repository.

# Supplementary Table S3. AIRRPrep and CLI-pRESTO: Comparison of Processing Results for the UMI-barcoded Illumina MiSeq 325+275 paired-end 5' RACE BCR mRNA validation workflow

The predefined UMI-barcoded 5' RACE MiSeq 325+275 workflow was run on the first 25,000 read pairs of SRR4026043. Step parameters are given in Supplementary Table S9.

| Step | Operation                | Stream  | Input           | AIRRPrep        | CLI-pRESTO      |
|------|--------------------------|---------|-----------------|-----------------|-----------------|
| 1–2  | FilterSeq.quality        | R1 / R2 | 25,000 / 25,000 | 22,285 / 20,424 | 22,285 / 20,424 |
| 3–4  | MaskPrimers.score        | R1 / R2 | 22,285 / 20,424 | 21,204 / 19,899 | 21,204 / 19,899 |
| 5    | PairSeq                  | paired  | 21,204 + 19,899 | 18,778 pairs    | 18,778 pairs    |
| 6–7  | BuildConsensus           | R1 / R2 | 18,778          | 16,022 / 16,039 | 16,022 / 16,039 |
| 8    | PairSeq                  | paired  | 16,022 + 16,039 | 16,020 pairs    | 16,020 pairs    |
| 9    | AssemblePairs.sequential | merged  | 16,020          | 15,493          | 15,493          |
| 10   | MaskPrimers.align        | merged  | 15,493          | 14,939          | 14,939          |
| 11   | ParseHeaders.collapse    | merged  | 14,939          | 14,939          | 14,939          |
| 12   | CollapseSeq              | merged  | 14,939          | 13,905          | 13,905          |
| 13   | SplitSeq.group           | merged  | 13,905          | 2,419           | 2,419           |

*Supplementary Table S3.* Number of records retained after each processing step of the UMI-barcoded 5' RACE MiSeq 325+275 workflow (SRR4026043, first 25,000 read pairs), for AIRRPrep and the equivalent command-line pRESTO implementation (records retained). At every step the two implementations wrote the same records: every sequence identifier was present on both sides with an identical nucleotide sequence, quality string and annotation set (compared as a multiset of field–value pairs, so that a repeated field counts as a difference while a reordering does not), with no missing, additional or altered record; PairSeq steps were compared mate by mate and the final table row by row. Record order differs, because command-line pRESTO writes records in the order its worker processes complete them. The comparison is produced by `compare_records.py` in the benchmark harness. Step parameters are given in Supplementary Table S9.

## Supplementary Table S4. AIRRPrep and CLI-pRESTO: Comparison of Processing Results for the Illumina MiSeq 2×250 paired-end BCR mRNA preprocessing workflow

The predefined Illumina MiSeq 2×250 BCR mRNA workflow was run on the first 25,000 read pairs of sample Replicate-1-M1 (accession: ERR346600). Step parameters are given in Supplementary Table S9.

| Step | Operation           | Stream  | AIRRPrep | CLI-pRESTO |
|------|---------------------|---------|----------|------------|
| 1    | AssemblePairs.align | R1 + R2 | 24,363   | 24,363     |
| 2    | FilterSeq.quality   | merged  | 24,361   | 24,361     |
| 3    | MaskPrimers.score   | merged  | 23,968   | 23,968     |
| 4    | MaskPrimers.score   | merged  | 23,573   | 23,573     |
| 5    | CollapseSeq         | merged  | 13,079   | 13,079     |
| 6    | SplitSeq.group      | merged  | 1,044    | 1,044      |

*Supplementary Table S4.* Number of records retained after each processing step of the Illumina MiSeq 2×250 paired-end BCR mRNA workflow (ERR346600, sample Replicate-1-M1, first 25,000 read pairs), for AIRRPrep and the equivalent command-line pRESTO implementation (records retained). At every step the two implementations wrote the same records (identical identifiers, sequences, quality strings and annotations; record order aside), as described for Supplementary Table S3. Step parameters are given in Supplementary Table S9.

## Supplementary Table S5. AIRRPrep and CLI-pRESTO: Comparison of Processing Results for the UMI-barcoded Illumina MiSeq 2×250 paired-end 5' RACE BCR mRNA validation workflow

The predefined UMI-barcoded 5' RACE MiSeq 2×250 workflow was run on the first 25,000 read pairs of SRR1383456. Step parameters are given in Supplementary Table S9.

| Step | Operation             | Stream  | AIRRPrep        | CLI-pRESTO      |
|------|-----------------------|---------|-----------------|-----------------|
| 1–2  | FilterSeq.quality     | R1 / R2 | 22,822 / 19,910 | 22,822 / 19,910 |
| 3–4  | MaskPrimers.score     | R1 / R2 | 20,421 / 19,049 | 20,421 / 19,049 |
| 5    | PairSeq               | R1 + R2 | 17,142 pairs    | 17,142 pairs    |
| 6–7  | BuildConsensus        | R1 / R2 | 4,097 / 4,089   | 4,097 / 4,089   |
| 8    | PairSeq               | R1 + R2 | 4,087 pairs     | 4,087 pairs     |
| 9    | AssemblePairs.align   | paired  | 1,888           | 1,888           |
| 10   | ParseHeaders.collapse | merged  | 1,888           | 1,888           |
| 11   | CollapseSeq           | merged  | 1,137           | 1,137           |
| 12   | SplitSeq.group        | merged  | 645             | 645             |

*Supplementary Table S5.* Number of records retained after each processing step of the UMI-barcoded 5' RACE MiSeq 2×250 workflow (SRR1383456, first 25,000 read pairs), for AIRRPrep and the equivalent command-line pRESTO implementation (records retained). At every step the two implementations wrote the same records (identical identifiers, sequences, quality strings and annotations; record order aside), as described for Supplementary Table S3. Step parameters are given in Supplementary Table S9.

## Supplementary Table S6. Validation of AIRRPrep on Cell Ranger-derived single-cell input

The Cell Ranger adapters were run on three Cell Ranger representations of the 10x Genomics sc5p\_v2\_hs\_PBMC\_10k human PBMC T-cell V(D)J dataset, exactly as a job runs them (parse, record validation, writing of `final.fasta`, `metadata.tsv` and `source_annotations.tsv`), with the opt-in for `productive` enabled so that it could be compared. The written files were then read back and compared record by record, keyed on `sequence_id`, with the source files parsed independently of the adapter code. This tests the fidelity of the parser, not the biological accuracy of Cell Ranger's reconstruction. The comparison is produced by `paper/test/validate_singlecell.py`, which also writes the per-field counts summarized below to `validation_report.tsv` and `validation_report.json`.

The five input files are the Cell Ranger 5.0.0 outputs published by 10x Genomics at [https://cf.10xgenomics.com/samples/cell-vdj/5.0.0/sc5p\\_v2\\_hs\\_PBMC\\_10k/](https://cf.10xgenomics.com/samples/cell-vdj/5.0.0/sc5p_v2_hs_PBMC_10k/) (prefix `sc5p_v2_hs_PBMC_10k_t_`); the MD5 checksums of the four single-part files equal those served by 10x, and the AIRR TSV matches in size.

|                                                                | All contigs | Filtered contigs | AIRR TSV           |
|----------------------------------------------------------------|-------------|------------------|--------------------|
| Source records                                                 | 15,377      | 10,817           | 10,817             |
| Output records                                                 | 15,377      | 10,817           | 10,817             |
| Distinct cell barcodes in output                               | 6,864       | 5,328            | 5,328              |
| Matched by <code>sequence_id</code>                            | 15,377      | 10,817           | 10,817             |
| Missing / additional records                                   | 0 / 0       | 0 / 0            | 0 / 0              |
| Duplicate output identifiers                                   | 0           | 0                | 0                  |
| Altered nucleotide sequences                                   | 0           | 0                | 0                  |
| Cell-barcode mismatches                                        | 0           | 0                | 0                  |
| <code>umi_count</code> / <code>read_count</code> mismatches    | 0 / 0       | 0 / 0            | 0 / 0              |
| <code>is_cell</code> / <code>high_confidence</code> mismatches | 0 / 0       | 0 / 0            | 0 / - <sup>a</sup> |
| <code>productive</code> mismatches (opt-in)                    | 0           | 0                | 0                  |
| Upstream annotation columns in output                          | 0           | 0                | 0                  |
| Rows in <code>source_annotations.tsv</code>                    | 15,377      | 10,817           | 10,817             |
| Source columns preserved there                                 | 18 of 18    | 18 of 18         | 32 of 32           |
| Rows that do not join back <sup>b</sup>                        | 0           | 0                | 0                  |
| Rows skipped by the adapter                                    | 0           | 0                | 0                  |

| Input SHA-256<br>(first 16 hex digits) | All contigs      | Filtered contigs | AIRR TSV         |
|----------------------------------------|------------------|------------------|------------------|
| FASTA                                  | 6c6d8404c797d4c4 | ca3a7412cfd84688 | —                |
| Annotations CSV                        | 200fc94628622bce | cf6b573fc00a52b0 | —                |
| AIRR TSV                               | —                | —                | 70376051c9e6cf2d |

*Supplementary Table S6.* Record-level comparison of AIRRPrep’s harmonized single-cell output with three Cell Ranger representations of the `sc5p_v2_hs_PBM10k` T-cell V(D)J dataset (`*_all_contig.fasta` with `*_all_contig_annotations.csv`; `*_filtered_contig.fasta` with `*_filtered_contig_annotations.csv`; `*_airr_rearrangement.tsv`; all prefixed `sc5p_v2_hs_PBM10k_t_`). Cell Ranger’s AIRR export stores the UMI count in `duplicate_count` and the read count in `consensus_count`. <sup>a</sup> The AIRR TSV has no `high_confidence` column; the value is left empty in all 10,817 output rows rather than assumed. <sup>b</sup> The lossless sidecar was written for each of the three routes and checked as a join partner of the harmonized output: every output record has exactly one row keyed on the same `sequence_id`, no identifier is duplicated, no row is orphaned in either direction, and every column the source file declared is present — including the V(D)J calls and CDR3s the harmonized record omits (`chain`, `v_gene`, `d_gene`, `j_gene`, `c_gene`, `cdr3`, `cdr3_nt` for the contig routes; `v_call`, `d_call`, `j_call`, `c_call`, `junction`, `junction_aa`, the alignments and the CIGARs for the AIRR TSV) and a column that is empty in every source row. With the `productive` opt-in disabled, the same adapter left `productive` empty in all output rows of the harmonized table, while the sidecar still carried the source’s own value. Every output FASTA header consisted of the identifier alone; the filtered-contig and AIRR TSV routes wrote byte-identical `final.fasta` files.

## Supplementary Table S7. Validation of the remaining single-cell input specifications

**Generic receptor FASTA.** The generic FASTA adapter was run on `sc5p_v2_hs_PBMC_10k_t_filtered_contig.fasta` (10,817 records) with each of the two ways of recovering a cell barcode, and with two inputs that must be refused. For G1 the barcode mapping file was written from `filtered_contig_annotations.csv` (`sequence_id`, `cell_id` = `barcode`, `umi_count` = `umis`, `read_count` = `reads`); for G2 the barcode was taken from the header with the rule `^([ACGT]++-\d+)_contig`. Outputs were compared record by record with the source as for Supplementary Table S6.

| Case | Input                                     | Expected | Observed                                                                                                                      |
|------|-------------------------------------------|----------|-------------------------------------------------------------------------------------------------------------------------------|
| G1   | FASTA + barcode mapping file              | accepted | 10,817 records, 5,328 cells; 0 missing, additional or altered records; 0 barcode, UMI-count or read-count mismatches          |
| G2   | FASTA + header rule                       | accepted | 10,817 records, 5,328 cells; 0 missing, additional or altered records; 0 barcode mismatches; counts empty (not in the source) |
| G3   | FASTA only, no barcode rule               | refused  | refused: the adapter requires a mapping file or header rule                                                                   |
| G4   | FASTA + mapping file missing one sequence | refused  | refused, naming the unassigned sequence (1 of 10,817)                                                                         |

A sequence is never given its own identifier as a cell identifier: that would turn N sequences into N artificial cells. In G1 and G2 `is_cell` and `high_confidence` are left empty, because a FASTA file carries no cell calling. Both accepted cases also wrote the lossless sidecar — the mapping file’s own columns in G1, the FASTA header in G2 — with 10,817 rows joining back to the 10,817 harmonized records without duplication or loss.

**TRUST4-based routes.** Precomputed TRUST4 output (`*_annot.fa` with `*_barcode_report.tsv`), raw 10x FASTQ and a 10x BAM share one parser. It assigns a contig to a cell only when TRUST4’s barcode report lists that contig as a chain of the cell (after an assembly run by AIRRPrep, the `cell_id` column of TRUST4’s `*_barcode_airr.tsv` is read as a second source and must agree), locating each contig by its identifier rather than by a fixed column position. The other contigs in `*_annot.fa`, which TRUST4 does not report as any cell’s chain, are excluded and counted; a contig claimed by two barcodes stops the run, and a file pair in which no contig maps to a barcode is refused. Raw reads and BAM files are assembled only in TRUST4’s barcode-aware mode, and a BAM whose CB/UB tags TRUST4 cannot use is refused rather than assembled without barcodes. Each run records these counts, and for BAM input extracted with samtools the number of reads carrying CB and UB tags, in its job record and `provenance.json`.

**Real-data evaluation.** The TRUST4 routes were run on the 10x Genomics `sc5p_v2_hs_PBMC_1k` human B-cell V(D)J library (5’ v2 chemistry), for which Cell Ranger 5.0.0 calls 83 cells with 166 high-confidence contigs (IGH 82, IGK 55, IGL 29) – the reference against which each run was compared. T2 used the first 598,163 read pairs of lane L001, with R1 read as a 16-nt barcode and 10-nt UMI (the 10x v2 preset); T1 uploaded the `annot.fa` and barcode report written by T2 as precomputed TRUST4 output; T3 used the library’s Cell Ranger `all_contig.bam` through the samtools-extraction route. T4 used the genome-aligned BAM of the 5’ gene-expression library captured from the same cells (Cell Ranger `possorted_genome_bam.bam`, 5.6 GB), through TRUST4’s own genome-coordinate BAM reader with the CB/UB tags; only the reads aligned

within the 24 IG/TCR regions spanned by the 620 genes of TRUST4’s `hg38_bcrtcr.fa` ( $\pm 10$  kb) were retrieved from the indexed public file (330,864 reads, 320,907 carrying CB), so unaligned reads, which TRUST4 also uses when given a whole BAM, are absent. The runs used AIRRPrep’s unmodified single-cell code in a Linux container with TRUST4 1.1.9 (bioconda build v1.1.9-r619) and its v1.1.9 human references (`hg38_bcrtcr.fa`, `human_IMGT+C.fa`), samtools 1.23.1 and eight threads (`paper/test/validate_trust4.py`, `paper/test/trust4-test`); the three accepted cases gave identical counts in three repeated runs. Datasets: Cell Ranger 5.0.0 outputs at [https://cf.10xgenomics.com/samples/cell-vdj/5.0.0/sc5p\\_v2\\_hs\\_PBMC\\_1k/](https://cf.10xgenomics.com/samples/cell-vdj/5.0.0/sc5p_v2_hs_PBMC_1k/).

| Case | Input                                                                                                  | Expected | Observed                                                                                                                                                                            |
|------|--------------------------------------------------------------------------------------------------------|----------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| T2   | 10x FASTQ, R1 + R2, 598,163 read pairs                                                                 | accepted | 9,240 contigs assembled; 627 reported as chains of 452 cells and written, 8,613 excluded; all 83 Cell Ranger cells present; CDR3 identical for 165/165 Cell Ranger cell-locus calls |
| T1   | T2’s <code>trust4_annot.fa</code> + <code>trust4_barcode_report.tsv</code>                             | accepted | 627 records, 452 cells; 627/627 records identical to T2 (identifier, cell, sequence)                                                                                                |
| T3   | <code>all_contig.bam</code> , samtools extraction                                                      | accepted | 2,833,786 single-end reads, all carrying CB and UB; 20,472 contigs; 622 records, 397 cells; all 83 Cell Ranger cells present; CDR3 identical for 165/165                            |
| T4   | 5’ gene-expression <code>possorted_genome_bam.bam</code> , IG/TCR-locus reads, genome-coordinate route | accepted | 6,461 contigs; 337 records, 279 cells (IG and TR chains); 79 of 83 Cell Ranger B cells present; CDR3 identical for 95/165 (light chains 77/83, heavy chains 17/81)                  |
| T3n1 | <code>all_contig.bam</code> , genome-coordinate route                                                  | refused  | refused before assembly: BAM aligned to contigs, not to a genome                                                                                                                    |
| T3n2 | TRUST4 <code>example.bam</code> (no CB/UB), genome-coordinate route                                    | refused  | refused: no contigs assembled with CB/UB tags; not retried without barcodes                                                                                                         |
| T3n3 | TRUST4 <code>example.bam</code> , samtools extraction                                                  | refused  | refused: none of 198 reads carries a CB tag                                                                                                                                         |

In T2, T3 and T4 no output `cell_id` equalled a contig identifier; the `cell_id` of TRUST4’s AIRR table agreed with the barcode report for all 562 (T2), 482 (T3) and 322 (T4) contigs it lists, with no conflicting assignment. For every (cell, locus) pair called by Cell Ranger, TRUST4 reported the same CDR3 nucleotide sequence for that barcode and locus (165/165 in T2 and in T3); heavy- and light-chain CDR3s were both identical in all 81 cells for which Cell Ranger calls both chains. TRUST4 called more cells than Cell Ranger (452 in T2, 397 in T3), which applies its own cell calling; seven T2 barcodes contain an undetermined base (N), because no 10x barcode whitelist was supplied. T4 establishes that the genome-coordinate route assembles and assigns cells correctly from a real 10x BAM, but its agreement with Cell Ranger’s V(D)J calls is limited by the input rather than by AIRRPrep: a gene-expression library is not enriched for receptor transcripts, its reads cover the heavy-chain CDR3 far more sparsely than the light chain’s, and the unaligned reads were not retrieved. It also calls T cells, which the gene-expression library contains and the B-cell V(D)J library does not.

Two defects were found by this evaluation and corrected before the runs reported above. First, the barcode and UMI positions of R1 were passed to TRUST4 with its deprecated `--barcodeRange/--umiRange` options, with which TRUST4 1.1.9 extracted an empty barcode for every read; the parser refused the resulting output (0 of 3,066 contigs assignable to a cell), and the positions are now passed with `--readFormat`. With the deprecated options TRUST4’s

`fastq-extractor` also terminated with a segmentation fault on some runs; this was not observed with `--readFormat`. Second, the samtools route kept only paired reads, whereas 10x BAM files store only the cDNA read, so it produced no input; it now assembles single-end reads. Both corrections are covered by automated tests.

An earlier draft reported these routes as validated on the example files distributed with TRUST4 (`example.bam`, `example_1.fq`, `example_2.fq`). Those are simulated bulk reads without cell barcodes, so that evaluation established only that TRUST4 was installed and executed, and it is replaced by the evaluation above.

*Supplementary Table S7.* Validation of the generic FASTA and TRUST4-based single-cell input specifications. G1–G4 were run on the `sc5p_v2_hs_PBM10k` T-cell filtered contigs (FASTA SHA-256 `ca3a7412cfd84688...`) by `paper/test/validate_singlecell.py`. T1–T4 were run on the `sc5p_v2_hs_PBM1k` library by `paper/test/validate_trust4.py`; input SHA-256 (first 16 hex digits): R1 FASTQ `8288981b3c56ef91`, R2 FASTQ `21431f7307eed796`, `all_contig.bam` `63a4051b0d40d60e`, IG/TCR-locus BAM `bcc2c1fdc3e3224a`, TRUST4 `example.bam` `73695c14b2e4f9ca`; the TRUST4 outputs uploaded in T1 are archived with the scripts (`trust4_annot.fa` `7ee564bce3d2bfa7`). Smoke tests (successful execution), parser-fidelity tests (record-level agreement with the source) and biological-accuracy tests (agreement of a reconstruction with an independent caller) are distinguished: G1–G4 and T1 are parser-fidelity tests, and T2–T4 compare TRUST4’s reconstruction, as parsed by AIRRPrep, with Cell Ranger’s calls on the same cells.

## Supplementary Table S8. Execution time of the predefined bulk workflows, overall and by processing step

Execution time was compared between AIRRPrep and the equivalent command-line pRESTO invocations for the three predefined bulk workflows, using the same inputs and the same parameters as in Supplementary Tables S3–S5. Both implementations were run on the same machine, each workflow three times per implementation, and the reported values are wall-clock times for the complete workflow, excluding upload and download.

**Scope of this comparison.** Everything reported in this table was measured on one Windows 11 workstation, on 25,000-read-pair subsets, and quantifies per-step execution overhead at that input size on that host. It is not a measurement of production-scale throughput, and it is not evidence of a platform-independent advantage: a substantial part of the margin outside read-pair assembly is the Windows process start-up cost described below, which a Linux host avoids by forking. A resource-matched Linux comparison at production input sizes has not been run, and no claim that depends on one is made anywhere in this work.

**Execution environment and CPU settings.** Both implementations ran natively on Windows 11 (no WSL, virtual machine or container), on one four-core, eight-thread CPU, from the same Python interpreter and pRESTO installation (Supplementary Table S10). The command-line scripts in Supplementary Table S9 do not set `--nproc`; pRESTO’s default for that option is the number of logical processors reported by the operating system (`multiprocessing.cpu_count()`), which was eight, so every parallel-capable command-line step ran with eight processes. AIRRPrep ran as one Celery 5.6.3 worker using the `solo` pool (one worker process, one task at a time; Celery does not support its default pool on Windows), started with `python -m celery -A app.workers.celery_worker.celery worker --pool=solo`, with a local Redis instance as broker and result backend. The worker was started before timing began and served all repetitions. Its CPU budget was eight logical processors, but only read-pair assembly (`AssemblePairs.py`), `AlignSets.py` and the external clustering tools of `ClusterSets.py` use more than one process in AIRRPrep; every other operation, including quality filtering, primer masking and consensus building, ran inside the single worker process. AIRRPrep therefore never used more processes than the command line for the same step. A command-line step on Windows starts each of its eight pRESTO worker processes as a fresh interpreter (the `spawn` start method), which Linux avoids by forking; that start-up cost is part of the per-step overhead measured here, so the margin outside assembly should be expected to be smaller on a Linux host. The benchmark harness (`benchmark.py`), the exact command-line scripts, the environment capture and all replicate times (`results.json`) are released with the source code, so the same comparison can be repeated on another host.

Two properties of the execution layer account for the difference. First, AIRRPrep invokes the pRESTO Python interface directly from a persistent worker process rather than launching a separate command-line process for each step, which removes the per-step interpreter start-up and module-import overhead incurred by a shell script that calls one pRESTO executable per operation. Measured step by step, the operations outside read-pair assembly completed about three times faster in aggregate on this basis alone:  $3.3\times$ ,  $3.4\times$  and  $3.0\times$  for the three workflows in the order tabulated below.

Second, the steps that can use more than one core do so, to a width taken from deployment configuration rather than left at one. This is what governs the assembly steps, and it is not a matter of passing a thread count through: the external aligners are invoked by pRESTO once per read pair and are hard-coded to a single thread each (`blastn` and `ublast`), and MUSCLE, used once per barcode set, exposes no thread option at all. In each case the machine can only be used by running

several of those invocations concurrently, which is what the command-line `--nproc` setting does and what AIRRPrep does here. The clustering tools invoked by `ClusterSets.py` – USEARCH, VSEARCH and CD-HIT-EST – instead accept a thread count directly and are given the same value.

Reference-guided assembly remains the most expensive step in absolute terms and is where the two implementations come closest: 236 s in AIRRPrep against 244 s on the command line for the 5' RACE 325+275 workflow, a factor of 1.03. That single step accounts for 68% of the command-line runtime of that workflow, which is why it shows the smallest overall margin. De novo assembly, which does the same per-pair work in Python without an external process, was faster in AIRRPrep in both workflows that use it: 68 s against 95 s on 25,000 pairs (1.40 $\times$ ) and 16 s against 35 s on 4,087 pairs (2.22 $\times$ ).

| Workflow                                               | AIRRPrep (s)     | CLI-pRESTO (s)  | Ratio         |
|--------------------------------------------------------|------------------|-----------------|---------------|
| UMI-barcoded 5' RACE MiSeq 325+275 (SRR4026043)        | 287.2 $\pm$ 17.6 | 346.1 $\pm$ 8.5 | 1.21 $\times$ |
| MiSeq 2 $\times$ 250 BCR mRNA (ERR346600)              | 80.1 $\pm$ 3.7   | 142.9 $\pm$ 1.0 | 1.78 $\times$ |
| UMI-barcoded 5' RACE MiSeq 2 $\times$ 250 (SRR1383456) | 41.0 $\pm$ 1.2   | 115.2 $\pm$ 0.3 | 2.81 $\times$ |

*Supplementary Table S8a.* Wall-clock execution time of the three predefined bulk workflows in AIRRPrep and in the equivalent command-line pRESTO implementation, measured on the inputs described in Supplementary Tables S3–S5. Values are mean  $\pm$  standard deviation over three replicates per implementation, and ratio is mean CLI-pRESTO time divided by mean AIRRPrep time, so values above one favour AIRRPrep. Replicate-to-replicate variation on the command-line side, which repeats a fully deterministic sequence of operations, was 0.3–2.5% of the mean. Both implementations ran on the same four-core, eight-thread machine, from the same Python 3.11.8 interpreter and the same pRESTO 0.7.9 installation, with command-line pRESTO at its default of eight processes per parallel step and AIRRPrep given a budget of eight logical processors (see the environment paragraph above); times exclude upload and download. The full software and hardware environment is listed in Supplementary Table S10. The two implementations wrote the same records at every step of all three workflows (Supplementary Tables S3–S5). These ratios describe this host and this input size only.

In what follows, we present the same comparison broken down by individual processing step, treating each of the three workflows as a separate table so that a step can be located by its position in the corresponding workflow. Per-step wall-clock time for both implementations is reported from one replicate of each workflow. Steps follow the numbering of Supplementary Tables S3–S5, which report the records retained at each one and at which both implementations wrote identical records throughout; the final annotation-table export is not listed there and is included here so that the steps account for the whole runtime. Ratio is CLI-pRESTO time divided by AIRRPrep time, so values above one favour AIRRPrep. Totals are the sum of the steps of the single replicate shown and therefore differ from the three-replicate means reported in Supplementary Table S8a above, which are the values to cite for overall runtime.

Two steps leave no timestamp of their own and are recovered as the residual of the run, that is, the measured total of that replicate minus every step that could be timed: † the first command-line step, because the inputs are copied into the run directory and a copy preserves the source modification time, and ‡ the last AIRRPrep step, whose output is written to the job output directory rather than to a numbered step directory. Each residual absorbs any unattributed overhead and is therefore an upper bound on that step. Where the recovery can be checked it holds: in the two workflows whose first and second steps are the same operation on opposite mate streams, the recovered first step falls within 2% of the measured second step.

**Supplementary Table S8b. UMI-barcoded 5' RACE MiSeq 325+275 (SRR4026043), by step**

| Step                        | AIRRPrep (s) | CLI-pRESTO (s) | Ratio        |
|-----------------------------|--------------|----------------|--------------|
| 1. FilterSeq.quality R1     | 1.7          | 13.2†          | 7.76×        |
| 2. FilterSeq.quality R2     | 1.6          | 12.9           | 8.06×        |
| 3. MaskPrimers.score R1     | 3.1          | 13.6           | 4.39×        |
| 4. MaskPrimers.score R2     | 2.5          | 13.3           | 5.32×        |
| 5. PairSeq                  | 2.7          | 3.3            | 1.22×        |
| 6. BuildConsensus R1        | 6.8          | 14.9           | 2.19×        |
| 7. BuildConsensus R2        | 5.6          | 14.8           | 2.64×        |
| 8. PairSeq                  | 2.0          | 2.8            | 1.40×        |
| 9. AssemblePairs.sequential | 236.0        | 243.8          | 1.03×        |
| 10. MaskPrimers.align       | 3.8          | 13.9           | 3.66×        |
| 11. ParseHeaders.collapse   | 1.3          | 2.0            | 1.54×        |
| 12. CollapseSeq             | 1.8          | 3.8            | 2.11×        |
| 13. SplitSeq.group          | 1.0          | 3.0            | 3.00×        |
| 14. ParseHeaders.table      | 0.5‡         | 0.7            | 1.40×        |
| <i>Total</i>                | <i>270.4</i> | <i>356.0</i>   | <i>1.32×</i> |

*Supplementary Table S8b.* Per-step wall-clock execution time for the UMI-barcoded 5' RACE MiSeq 325+275 workflow (SRR4026043), for AIRRPrep and command-line pRESTO, from one replicate of each. Steps marked † or ‡ are recovered residuals rather than directly measured times, as explained above.

**Supplementary Table S8c. MiSeq 2×250 BCR mRNA (ERR346600), by step**

| Step                   | AIRRPrep (s) | CLI-pRESTO (s) | Ratio        |
|------------------------|--------------|----------------|--------------|
| 1. AssemblePairs.align | 67.8         | 94.8†          | 1.40×        |
| 2. FilterSeq.quality   | 2.1          | 12.2           | 5.81×        |
| 3. MaskPrimers.score V | 6.6          | 14.5           | 2.20×        |
| 4. MaskPrimers.score C | 2.5          | 14.3           | 5.72×        |
| 5. CollapseSeq         | 1.6          | 3.7            | 2.31×        |
| 6. SplitSeq.group      | 0.9          | 2.9            | 3.22×        |
| 7. ParseHeaders.table  | 0.5‡         | 0.7            | 1.40×        |
| <i>Total</i>           | <i>82.0</i>  | <i>143.1</i>   | <i>1.75×</i> |

*Supplementary Table S8c.* Per-step wall-clock execution time for the MiSeq 2×250 BCR mRNA workflow (ERR346600), for AIRRPrep and command-line pRESTO, from one replicate of each. Steps marked † or ‡ are recovered residuals rather than directly measured times, as explained above.

**Supplementary Table S8d. UMI-barcoded 5' RACE MiSeq 2×250 (SRR1383456), by step**

| Step                      | AIRRPrep (s) | CLI-pRESTO (s) | Ratio        |
|---------------------------|--------------|----------------|--------------|
| 1. FilterSeq.quality R1   | 1.5          | 11.0†          | 7.33×        |
| 2. FilterSeq.quality R2   | 1.5          | 11.1           | 7.40×        |
| 3. MaskPrimers.score R1   | 2.7          | 12.2           | 4.52×        |
| 4. MaskPrimers.score R2   | 8.9          | 12.6           | 1.42×        |
| 5. PairSeq                | 2.4          | 2.9            | 1.21×        |
| 6. BuildConsensus R1      | 4.1          | 9.9            | 2.41×        |
| 7. BuildConsensus R2      | 4.3          | 13.1           | 3.05×        |
| 8. PairSeq                | 0.5          | 1.4            | 2.80×        |
| 9. AssemblePairs.align    | 15.8         | 35.0           | 2.22×        |
| 10. ParseHeaders.collapse | 0.2          | 0.9            | 4.50×        |
| 11. CollapseSeq           | 0.2          | 2.0            | 10.00×       |
| 12. SplitSeq.group        | 0.1          | 2.1            | 21.00×       |
| 13. ParseHeaders.table    | 0.1‡         | 0.6            | 6.00×        |
| <i>Total</i>              | <i>42.3</i>  | <i>114.8</i>   | <i>2.71×</i> |

*Supplementary Table S8d.* Per-step wall-clock execution time for the UMI-barcoded 5' RACE MiSeq 2×250 workflow (SRR1383456), for AIRRPrep and command-line pRESTO, from one replicate of each. Steps marked † or ‡ are recovered residuals rather than directly measured times, as explained above.

## Supplementary Table S9. Parameters of the predefined bulk workflows and the reference command-line scripts

Each predefined workflow was run in AIRRPrep and, for comparison, as a shell script that invokes one pRESTO console program per operation – the form in which these protocols are ordinarily run. The two sides were given identical parameters; the scripts below are the reference implementations used for Supplementary Tables S3–S5 and S8 and are reproduced in full, so that both the parameter set and the comparison can be repeated exactly. Line breaks marked with a trailing backslash are shell continuations and are not part of the command.

The primer, template-switch and internal C-region FASTA files, and the IMGT human IGV reference used for reference-guided assembly, are the files distributed with pRESTO for these protocols. Inputs are the first 25,000 read pairs of each accession, as written by `fastq-dump --split-files`; their sequences are identical, record for record, to the first 25,000 records of the ENA FASTQ files of each run, and so are their quality strings except for SRR4026043, which ENA distributes with a constant placeholder quality (`verify_inputs_ena.py` in the benchmark harness); because the subsets are prefixes rather than random samples, they are reproducible without a seed.

### UMI-barcoded 5' RACE MiSeq 325+275 (SRR4026043)

```
FilterSeq.py quality -s SRR4026043_1.fastq -q 20 --outname HD09N-R1
FilterSeq.py quality -s SRR4026043_2.fastq -q 20 --outname HD09N-R2
MaskPrimers.py score -s HD09N-R1_quality-pass.fastq \
    -p AbSeq_R1_Human_IG_Primers.fasta --start 0 --mode cut --outname HD09N-R1
MaskPrimers.py score -s HD09N-R2_quality-pass.fastq \
    -p AbSeq_R2_TS.fasta --start 17 --barcode --mode cut --maxerror 0.5 \
    --outname HD09N-R2
PairSeq.py -1 HD09N-R1_primers-pass.fastq -2 HD09N-R2_primers-pass.fastq \
    --2f BARCODE --coord sra
BuildConsensus.py -s HD09N-R1_primers-pass_pair-pass.fastq --bf BARCODE \
    --pf PRIMER --prcons 0.6 --maxerror 0.1 --maxgap 0.5 --outname HD09N-R1
BuildConsensus.py -s HD09N-R2_primers-pass_pair-pass.fastq --bf BARCODE \
    --maxerror 0.1 --maxgap 0.5 --outname HD09N-R2
PairSeq.py -1 HD09N-R1_consensus-pass.fastq -2 HD09N-R2_consensus-pass.fastq \
    --coord presto
AssemblePairs.py sequential -1 HD09N-R2_consensus-pass_pair-pass.fastq \
    -2 HD09N-R1_consensus-pass_pair-pass.fastq -r IMGT_Human_IG_V.fasta \
    --coord presto --rc tail --scanrev --1f CONSCOUNT --2f CONSCOUNT PRCONS \
    --aligner blastn --outname HD09N-C
MaskPrimers.py align -s HD09N-C_assemble-pass.fastq \
    -p AbSeq_Human_IG_InternalCRegion.fasta --maxlen 100 --maxerror 0.3 \
    --mode tag --revpr --skiprc --pf CREGION --outname HD09N-C
ParseHeaders.py collapse -s HD09N-C_primers-pass.fastq -f CONSCOUNT --act min
CollapseSeq.py -s HD09N-C_primers-pass_reheader.fastq -n 20 --inner \
    --uf CREGION --cf CONSCOUNT --act sum --outname HD09N-C
SplitSeq.py group -s HD09N-C_collapse-unique.fastq -f CONSCOUNT --num 2 \
    --outname HD09N-C
ParseHeaders.py table -s HD09N-C_atleast-2.fastq -f ID CREGION CONSCOUNT DUPCOUNT
```

### MiSeq 2×250 BCR mRNA (ERR346600)

```
AssemblePairs.py align -1 ERR346600_2.fastq -2 ERR346600_1.fastq \
    --coord sra --rc tail --outname M1
```

```

FilterSeq.py quality -s M1_assemble-pass.fastq -q 20 --outname M1
MaskPrimers.py score -s M1_quality-pass.fastq -p Greiff2014_VPrimers.fasta \
    --start 4 --mode mask --pf VPRIMER --outname M1-FWD
MaskPrimers.py score -s M1-FWD_primers-pass.fastq -p Greiff2014_CPrimers.fasta \
    --start 4 --mode cut --revpr --pf CPRIMER --outname M1-REV
CollapseSeq.py -s M1-REV_primers-pass.fastq -n 20 --inner --uf CPRIMER \
    --cf VPRIMER --act set --outname M1
SplitSeq.py group -s M1_collapse-unique.fastq -f DUPCOUNT --num 2 --outname M1
ParseHeaders.py table -s M1_atleast-2.fastq -f ID DUPCOUNT CPRIMER VPRIMER

```

## **UMI-barcoded 5' RACE MiSeq 2×250 (SRR1383456)**

```

FilterSeq.py quality -s SRR1383456_1.fastq -q 20 --outname MS12_R1
FilterSeq.py quality -s SRR1383456_2.fastq -q 20 --outname MS12_R2
MaskPrimers.py score -s MS12_R1_quality-pass.fastq \
    -p Stern2014_CPrimers.fasta --start 15 --mode cut --barcode \
    --outname MS12_R1
MaskPrimers.py score -s MS12_R2_quality-pass.fastq \
    -p Stern2014_VPrimers.fasta --start 0 --mode mask --outname MS12_R2
PairSeq.py -1 MS12_R1_primers-pass.fastq -2 MS12_R2_primers-pass.fastq \
    --1f BARCODE --coord sra
BuildConsensus.py -s MS12_R1_primers-pass_pair-pass.fastq --bf BARCODE \
    --pf PRIMER --prcons 0.6 --maxerror 0.1 --maxgap 0.5 --outname MS12_R1
BuildConsensus.py -s MS12_R2_primers-pass_pair-pass.fastq --bf BARCODE \
    --maxerror 0.1 --maxgap 0.5 --outname MS12_R2
PairSeq.py -1 MS12_R1_consensus-pass.fastq -2 MS12_R2_consensus-pass.fastq \
    --coord presto
AssemblePairs.py align -1 MS12_R2_consensus-pass_pair-pass.fastq \
    -2 MS12_R1_consensus-pass_pair-pass.fastq --coord presto --rc tail \
    --1f CONSCOUNT --2f CONSCOUNT PRCONS --outname MS12
ParseHeaders.py collapse -s MS12_assemble-pass.fastq -f CONSCOUNT --act min
CollapseSeq.py -s MS12_assemble-pass_reheader.fastq -n 20 --inner --uf PRCONS \
    --cf CONSCOUNT --act sum --outname MS12
SplitSeq.py group -s MS12_collapse-unique.fastq -f CONSCOUNT --num 2 \
    --outname MS12
ParseHeaders.py table -s MS12_atleast-2.fastq -f ID PRCONS CONSCOUNT DUPCOUNT

```

*Supplementary Table S9.* Reference command-line implementations of the three predefined bulk workflows. The AIRRPrep runs compared against them in Supplementary Tables S3–S5 and S8 used the same operations, in the same order, with the same parameter values; each pRESTO program above corresponds to one component of the predefined workflow, and read-stream assignment (R1, R2, paired or merged) follows the `-s/-1/-2` arguments shown. Log options (`--log`) are omitted here for brevity and were present on both sides.

## Supplementary Table S10. Software and hardware environment

The runtime comparison in Supplementary Table S8 was executed with both implementations sharing one machine, one Python interpreter and one pRESTO installation, so that the measured difference reflects the execution layer rather than the environment. The TRUST4-based single-cell evaluation of Supplementary Table S7 was run in a separate Linux test container (middle part of the table) using the same TRUST4 release as the deployment image, and the lower part lists the versions pinned in the deployment image.

| Component                                             | Version / setting                                                                                 |
|-------------------------------------------------------|---------------------------------------------------------------------------------------------------|
| <i>Benchmark environment (Supplementary Table S8)</i> |                                                                                                   |
| Operating system                                      | Windows 11 (build 10.0.26200)                                                                     |
| CPU                                                   | Intel Core i7-10510U, 4 physical cores / 8 logical processors                                     |
| Memory                                                | 16 GB                                                                                             |
| Storage                                               | NVMe SSD                                                                                          |
| Execution                                             | native Windows; no WSL, virtual machine or container                                              |
| CLI-pRESTO parallelism                                | default <code>--nproc = 8</code> ( <code>multiprocessing.cpu_count()</code> )                     |
| AIRRPrep CPU budget                                   | 8 logical processors (whole machine); parallel only in AssemblePairs, AlignSets and ClusterSets   |
| Celery worker                                         | one worker, <code>--pool=solo</code> , started before timing                                      |
| Python                                                | 3.11.8                                                                                            |
| pRESTO                                                | 0.7.9                                                                                             |
| Biopython                                             | 1.87                                                                                              |
| NumPy / pandas / SciPy                                | 2.4.4 / 3.0.3 / 1.17.1                                                                            |
| FastAPI / Celery / uvicorn                            | 0.136.1 / 5.6.3 / 0.46.0                                                                          |
| Redis (client / broker)                               | redis-py 7.4.0 / Redis server 8.2.5, local                                                        |
| BLAST+ ( <code>blastn</code> )                        | 2.17.0+                                                                                           |
| USEARCH                                               | 7.0.1090 (component tests only; proprietary, not redistributed)                                   |
| VSEARCH                                               | 2.31.0                                                                                            |
| MUSCLE                                                | 3.8.31 (component tests)                                                                          |
| <i>Component matrix (Supplementary Table S2)</i>      |                                                                                                   |
| Host cases                                            | the benchmark environment above; native pRESTO from the same virtual environment AIRRPrep runs in |

| Component                                              | Version / setting                                                                                                                                                                            |
|--------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Container cases                                        | <code>paper/test/cluster-test/Dockerfile</code> on<br><code>python:3.12-slim</code> ; pRESTO 0.7.9, Biopython 1.87,<br>NumPy 2.4.4, SciPy 1.17.1, pandas 3.0.3                               |
| External tools in the container                        | VSEARCH 2.31.0 and BLAST+ 2.17.0, official builds<br>pinned by SHA-256 in <code>fetch_tools.sh</code>                                                                                        |
| <i>TRUST4 evaluation (Supplementary Table S7)</i>      |                                                                                                                                                                                              |
| Container                                              | Docker Desktop on Windows 11 (WSL 2), 8 CPUs, 7.6 GB;<br><code>paper/test/trust4-test/Dockerfile</code>                                                                                      |
| TRUST4                                                 | 1.1.9 (bioconda build, reports v1.1.9-r619), 8 threads                                                                                                                                       |
| TRUST4 references                                      | <code>hg38_bcrtcr.fa</code> , <code>human_IMGT+C.fa</code> from TRUST4 v1.1.9                                                                                                                |
| samtools                                               | 1.23.1 (htslib 1.23.1)                                                                                                                                                                       |
| Python                                                 | 3.12.14, with Biopython 1.87 and pydantic 2.13.4                                                                                                                                             |
| <i>Deployment container image (backend/Dockerfile)</i> |                                                                                                                                                                                              |
| Base image                                             | <code>python:3.12-slim</code>                                                                                                                                                                |
| Published image                                        | <code>ghcr.io/computimm/airrprep-backend@sha256:1cc4<br/>dcc98bb2869f9b6d6300887f6d593e1ba2e2464ae1e864<br/>113a62a6da4e96</code> (release v1.0.0)                                           |
| TRUST4                                                 | 1.1.9, built from source at the tagged release, with its<br>bundled reference set                                                                                                            |
| BLAST+, samtools, MUS-<br>CLE 3, VSEARCH, CD-HIT       | Debian packages; the version each build resolved is<br>recorded in the image as <code>/app/TOOL_VERSIONS.txt</code> , and<br>for the published image in <code>docker/tool-versions.md</code> |
| USEARCH                                                | not included                                                                                                                                                                                 |
| Default clustering tool / refer-<br>ence aligner       | VSEARCH / BLAST+ (both in the image); USEARCH<br>is never a default                                                                                                                          |
| Frontend                                               | React with TanStack Start                                                                                                                                                                    |

*Supplementary Table S10.* Software versions and hardware used for the reported measurements, and the versions pinned in the deployment image. USEARCH is proprietary; it is not required by any predefined workflow, is not selected by any default, and is not included in the image. It was installed locally only for the separate `ClusterSets` test described under Supplementary Table S2, and a deployment holding a USEARCH licence supplies its own binary (environment variable `USEARCH_BIN`). The Debian-packaged tools are not version-pinned in the Dockerfile; the versions an image contains are read back from the image itself.

## Supplementary Table S11. Comparison with related AIRR-seq tools

Table S11 compares AIRRPrep with the tools discussed in the Introduction, each in its own column. Every property is defined narrowly enough to be checked, and every cell about another tool is taken from that tool’s publication or official documentation (source given in brackets, listed below the table). “Not evaluated” means the authors found no statement on the property in those sources and did not test it; it does not mean the tool lacks the property. The comparison is not a benchmark of the other tools.

| Property                                                                           | AIRRPrep                                                                             | pRESTO                                                        | nf-core/airrflow                                                                                                           | VDJServer                                                                                                         | ARGalaxy                                                                                                             |
|------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| Interface                                                                          | Web application; container for self-hosting                                          | Command-line programs [1]                                     | Nextflow pipeline run from the command line [2]                                                                            | Web browser interface [3]                                                                                         | Galaxy web interface [4]                                                                                             |
| Stated scope                                                                       | Preprocessing of bulk reads; harmonization of single-cell inputs; no gene assignment | Preprocessing of raw repertoire reads [1]                     | Bulk and single-cell data from raw reads or assembled sequences, through annotation, clonal analysis and lineage trees [2] | Preprocessing and quality control, V(D)J assignment, repertoire characterization and comparison, data sharing [3] | Analysis and visualization of repertoire data from IMGT/HighV-QUEST output, incl. diversity, SHM, clonality, CSR [4] |
| User composition of individual pRESTO operations in any order                      | Yes, interactive builder over 47 operations                                          | Yes, by writing command scripts [1]                           | No; a defined pipeline configured by parameters and protocol profiles [2]                                                  | Not evaluated                                                                                                     | Not evaluated                                                                                                        |
| Check of a composed operation sequence against the supplied files before execution | Yes (Supplementary Note S1)                                                          | No; each command checks its own arguments when it runs [1]    | Not evaluated; the step sequence is defined by the pipeline [2]                                                            | Not evaluated                                                                                                     | Not evaluated                                                                                                        |
| Per-step record counts shown as a workflow graph                                   | Yes, exportable                                                                      | Per-step log files, tabulated by <code>ParseLog.py</code> [1] | Not evaluated                                                                                                              | Not evaluated                                                                                                     | Not evaluated                                                                                                        |
| Reusable, machine-readable workflow definition                                     | Yes, JSON export with companion-file checksums                                       | Shell scripts written by the user                             | Versioned pipeline with parameter files and profiles [2]                                                                   | Not evaluated                                                                                                     | Not evaluated                                                                                                        |

| Property                                    | AIRRPrep                                                                                        | pRESTO                                      | nf-core/airrflow                                            | VDJServer                               | ARGalaxy      |
|---------------------------------------------|-------------------------------------------------------------------------------------------------|---------------------------------------------|-------------------------------------------------------------|-----------------------------------------|---------------|
| Single-cell inputs                          | Seven input specifications harmonized to an internal 8-field record (Supplementary Table S1)    | Not a stated scope [1]                      | 10x raw reads and single-cell AIRR rearrangement tables [2] | Not evaluated                           | Not evaluated |
| Output standard                             | pRESTO files (bulk); internal schema, AIRR compliance not claimed (single-cell)                 | pRESTO-annotated FASTA/FASTQ and tables [1] | AIRR-compliant outputs [2]                                  | Repository in the AIRR Data Commons [3] | Not evaluated |
| Upstream annotations of a single-cell input | Excluded from the harmonized record; preserved unchanged in <code>source_annotations.tsv</code> | Not a stated scope [1]                      | Not evaluated                                               | Not evaluated                           | Not evaluated |

*Supplementary Table S11.* AIRRPrep compared with pRESTO, nf-core/airrflow, VDJServer and ARGalaxy. Every cell about another tool is taken from the source bracketed in that cell; the sources are listed in full below. Documentation was accessed in September 2026. AIRRPrep’s distinguishing combination is interactive composition of individual pRESTO operations with contract-based checking of the composition against the uploaded files before execution; the other tools cover scopes AIRRPrep does not, notably annotation and downstream analysis.

## References for Supplementary Table S11

1. Vander Heiden JA, Yaari G, Uduman M, Stern JNH, O’Connor KC, Hafler DA, Vigneault F, Kleinstein SH. pRESTO: a toolkit for processing high-throughput sequencing raw reads of lymphocyte receptor repertoires. *Bioinformatics*. 2014;30(13):1930–1932. <https://doi.org/10.1093/bioinformatics/btu138>. Accompanying documentation: pRESTO documentation, <https://presto.readthedocs.io> (accessed September 2026).
2. Gabernet G, Marquez S, Bjornson R, Peltzer A, Meng H, Aron E, Lee NY, Jensen CG, Ladd D, Polster M, Hanssen F, Heumos S, Yaari G, Kowarik MC, Nahnsen S, Kleinstein SH. nf-core/airrflow: an adaptive immune receptor repertoire analysis workflow employing the Immcantation framework. *PLoS Computational Biology*. 2024;20(7):e1012265. <https://doi.org/10.1371/journal.pcbi.1012265>. Accompanying documentation: nf-core/airrflow documentation, <https://nf-co.re/airrflow> (accessed September 2026).

3. Christley S, Scarborough W, Salinas E, Rounds WH, Toby IT, Fonner JM, Levin MK, Kim M, Mock SA, Jordan C, Ostmeyer J, Buntzman A, Rubelt F, Davila ML, Monson NL, Scheuermann RH, Cowell LG. VDJServer: a cloud-based analysis portal and data commons for immune repertoire sequences and rearrangements. *Frontiers in Immunology*. 2018;9:976. <https://doi.org/10.3389/fimmu.2018.00976>. Accompanying documentation: VDJServer documentation, <https://vdjserver.readthedocs.io> (accessed September 2026).
4. IJspeert H, van Schouwenburg PA, van Zessen D, Pico-Knijnenburg I, Stubbs AP, van der Burg M. Antigen receptor Galaxy: a user-friendly, web-based tool for analysis and visualization of T and B cell receptor repertoire data. *The Journal of Immunology*. 2017;198(10):4156–4165. <https://doi.org/10.4049/jimmunol.1601921>.